

CS-377 Parallel Programming

Final Assignment

Eliot Cowley

May 16, 2015

1 The S'mores Problem

Before starting this problem, my first decision was whether to use Ruby/Rinda or Java. I decided to go with Java since I was more familiar with it and preferred it in general. Having only done the one Producer/Consumer problem in Rinda, I was still a little hazy on how it all worked, and was more comfortable with synchronized methods in Java.

As with most assignments, I found it helpful to first sketch out how I was going to structure my program before doing any actual coding. At first, I was confused as to why I needed four threads—couldn't I just have three threads representing the children waiting to eat, with one synchronizing monitor? But then I realized that I needed a **Chooser** thread that would be waiting to choose ingredients and signaling the correct **Child**.

So I organized my program into four classes: the main **S'mores** class that would initialize the threads and monitor; the **Child** class, which would try to pick up ingredients and make s'mores; the **Chooser** class, which would choose ingredients to put down; and the **Monitor** class, which would synchronize the shared variables and make sure that deadlock and race conditions were avoided.

Having decided on the structure of my program, I began adding fields and methods to my classes. **Child** would have an index between 0 and 2 (I used constants so the program would be scalable), a reference to the **Monitor**, and an **int** representing the number of s'mores it has eaten. As for methods, it would have a **pickUpIngredients** method that would go through the monitor to try to pick up the ingredients off the table. It would also have a **run** method that would be that thread's lifecycle. Finally, it would have a **delay** method that would make the thread sleep for a random number of milliseconds determined by another method, **randomInt**. A **Child**'s lifecycle would consist of trying to pick up ingredients, making a s'more, delaying for a few milliseconds, and then repeating.

The **Chooser** class's fields would include a reference to the **Monitor** and an **int** representing the index of the **Child** he has chosen that turn. It would have a **run** method, as well as **putDownIngredients** and **choose** methods, which would go through the **Monitor**. The life cycle of the **Chooser** would consist of choosing a **Child** who would eat a s'more that turn and putting down the correct ingredients on the table.

The **Monitor** class would have several fields, including: an **int** representing the index of the child who will be making a s'more that turn; a **boolean** representing whether there are ingredients currently on the table; an array of **ints** representing the ingredients on the table; and a constant representing the number of children who will be eating s'mores in order to make the program scalable. It would have a synchronized method **putDownIngredients** which would take the index of the eater as a parameter, and would wait for there to be no ingredients on the table before inserting ingredients into the array, putting down only the ingredients of the two non-eater children. It would also have a synchronized method **pickUpIngredients** which would also take an eater as a parameter, however this method would wait both for there to be ingredients on the table and for the current eater to be equal to the caller thread. There would also be a synchronized **choose** method which randomly chose a number between 0 and 2 and assigned that index as

the eater. Finally there would be a `getName` method that took an integer as input and returned a `String` representing the ingredient of that index. The `Child` and `Chooser` threads would all call their methods through this `Monitor` to avoid deadlock and race conditions.

My solution works because any time a thread needs to access or modify a shared variable, it does so through synchronized methods in a monitor. This way, no two threads are stuck fighting for the same resource. I cannot say with certainty that this solution is fair, however after extensive testing I have found that each child generally eats the same amount of s'mores and no one child frequently eats the most.

2 The Water Molecule Problem

For this problem, I decided that to simplify things, the values that would be passed along channels would be integers. I declared 0 to be `hydrogen`, 1 `oxygen`, and 2 `water`. I made the two `hydrogen` processes and one `oxygen` process continuously output their respective data to both the `water` process and `print.stream` process. The `water` process waits until it has two `hydrogens` and an `oxygen` before it passes a `water` to `print.stream`. `print.stream` alternates over all four processes, choosing them arbitrarily and writing to the screen when something is produced. It then delays for a certain amount of milliseconds before continuing, so the user can more easily read the output.

The program is run in parallel wherever possible. The `hydrogen` and `oxygen` processes output to the `water` and `print.stream` processes in parallel. The `water` process reads in `hydrogen` and `oxygen` in parallel. Additionally, the main process `waterMolecule` runs the `hydrogen`, `oxygen`, `water`, and `print.stream` processes in parallel. There were a couple of blocks that I couldn't run in parallel, however—for instance, `water` can only output to `print.stream` once it has read in two `hydrogens` and an `oxygen`, so the output statement was done sequentially; `print.stream` chooses processes in an `ALT` block and then pauses afterwards, which is also done sequentially. These sequential blocks were necessary to ensure the program's accuracy.

There are a couple of things about the output of `print.stream` that seem odd, but make sense. First, sometimes `water` molecules are printed out before two `hydrogens` and an `oxygen` have been. However, this does not mean that the atoms required for a `water` molecule have not been created yet. In fact, they may have already been created but are just waiting on `ALT` to choose them to print to the screen. Because there is a delay after each print, there is time for `water` molecules to be created, so `ALT` may arbitrarily choose to print out a `water` molecule that is waiting before it prints out the atoms. I proved this by running my program without the pause after printing, and it always printed out at least two `hydrogens` and one `oxygen` before printing out a `water`. If I wanted to make for more accurate printing, I could have put a delay in the `water` process to make sure that it gets formed and then printed after the atoms had.

Second, the output always appears to be the same, even though technically `print.stream` is choosing processes "arbitrarily" with `ALT`. There is always a start-up section of output that is different from the rest, which then settles into a pattern of `water`, `oxygen`, `hydrogen`, `hydrogen`. I got a similar result, albeit with different patterns, when I switched the order of the guards. Additionally, when I tried using `PRI ALT`, it eventually settled into choosing processes from highest to lowest priority. The reason for `ALT` choosing processes with seeming regularity is likely because it is being fair, giving each process a chance to execute in order to avoid starvation. Therefore, `ALT` is not *completely* arbitrary, though it never gives unequal preference to any one process. Deadlock is avoided because there is no contention for resources; the `hydrogen` and `oxygen` processes produce integers for both `water` and `print.stream` equally. For the above reasons, my solution is correct.